

# 00\_fundamentals\_skimage

September 11, 2019

To participate, you'll need to git clone (or download the .zip from GitHub):

<https://github.com/imagexd/2019-tutorial-skimage>

You can do that in git using:

If you have already cloned the material, please issue `git pull` now and reload the notebook to ensure that you have the latest updates.

```
[1]: %matplotlib inline
```

- <https://scikit-image.org>
- Take a look at the gallery, API docs
- The underlying stack: NumPy, SciPy (especially `scipy.ndimage`), Cython

But, before we can start using `scikit-image`, we need to understand how images are represented. For this, we will use NumPy.

## 0.0.1 A quick 5-minute quiz, to refresh our NumPy skills!

Start with:

```
import numpy as np
```

Then answer:

1. What version of NumPy do you have installed? (Hint: `np.__version__`)
2. How do I construct a 300x500 array of zeros?
3. How do I generate an array with the numbers from 0 to 9?
4. How do I take the square of this array?
5. What is the data-type of this array? (Hint: `x.dtype`) What is the data-type of the array `[1, 2, 3]`?
6. What is the shape of the following array: `[[1, 2, 3], [4, 5, 6]]`
7. What, do you predict, is the result of `[[1, 2, 3], [4, 5, 6]] + [10, 11, 12]`? (Hint: broadcasting)
8. What, do you predict, is the result of `np.array([1,2,3], dtype=np.uint8) - 2`? (Hint: the 'u' stands for unsigned)
9. What, do you predict, is the result of `x = np.ones((5, 1)); x @ x.T` (Hint: @ is matrix multiplication)
10. What, do you predict, is the result of `x = np.array([1, 2, 3]); x[x < 3]` (Hint: what is `x > 3`?)

# 1 Images are numpy arrays


Images are represented in `scikit-image` using standard `numpy` arrays. This allows maximum inter-operability with other libraries in the scientific Python ecosystem, such as `matplotlib` and `scipy`.

Let's see how to build a grayscale image as a 2D array:

```
[2]: import numpy as np
      from matplotlib import pyplot as plt
```

```
[3]: random_image = np.random.random([500, 500])

      plt.imshow(random_image, cmap='gray')
      plt.colorbar();
```



The same holds for “real-world” images:


```
[4]: from skimage import data

      coins = data.coins()

      print('Type:', type(coins))
      print('dtype:', coins.dtype)
      print('shape:', coins.shape)
```

```
plt.imshow(coins, cmap='gray');
```


```
Type: <class 'numpy.ndarray'>  
dtype: uint8  
shape: (303, 384)
```



A color image is a 3D array, where the last dimension has size 3 and represents the red, green, and blue channels:


```
[5]: cat = data.chelsea()  
  
print("Shape:", cat.shape)  
print("Values min/max:", cat.min(), cat.max())  
  
plt.imshow(cat);
```

```
Shape: (300, 451, 3)  
Values min/max: 0 231
```



These are *just NumPy arrays*. E.g., we can make a red square by using standard array slicing and manipulation:

```
[6]: cat[10:110, 10:110, :] = [255, 0, 0] # [red, green, blue]
plt.imshow(cat);
```



Images can also include transparent regions by adding a 4th dimension, called an *alpha layer*.

### 1.0.1 Other shapes, and their meanings

Image type	Coordinates
2D grayscale	(row, column)
2D multichannel	(row, column, channel)
3D grayscale (or volumetric)	(plane, row, column)
3D multichannel	(plane, row, column, channel)

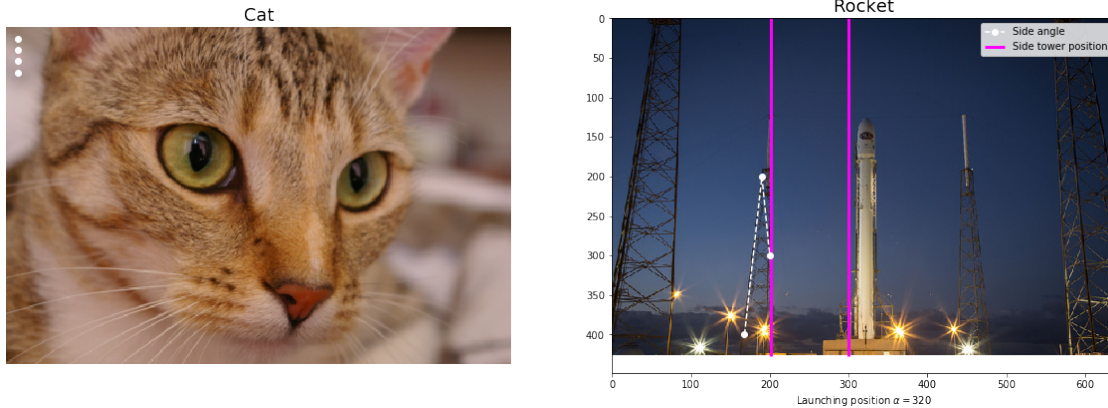
## 1.1 Displaying images using matplotlib

```
[7]: from skimage import data
```

```
img0 = data.chelsea()  
img1 = data.rocket()
```

```
[8]: import matplotlib.pyplot as plt
```

```
f, (ax0, ax1) = plt.subplots(1, 2, figsize=(20, 10))  
  
ax0.imshow(img0)  
ax0.set_title('Cat', fontsize=18)  
ax0.scatter([10, 10, 10, 10], [10, 20, 30, 40], color='white')  
ax0.axis('off')  
  
ax1.imshow(img1)  
ax1.set_title('Rocket', fontsize=18)  
ax1.set_xlabel(r'Launching position  $\alpha=320^\circ$ )  
  
ax1.vlines([202, 300], 0, img1.shape[0], colors='magenta',  
           linewidth=3, label='Side tower position')  
ax1.plot([168, 190, 200], [400, 200, 300], color='white',  
         linestyle='--', marker='o', label='Side angle')  
  
ax1.legend();
```



For more on plotting, see the [Matplotlib documentation](#) and [pyplot API](#).

## 1.2 Data types and image values

In literature, one finds different conventions for representing image values:

0 - 255    where 0 is black, 255 is white

0 - 1      where 0 is black, 1 is white

scikit-image supports both conventions—the choice is determined by the data-type of the array.

E.g., here, I generate two valid images:


```
[9]: linear0 = np.linspace(0, 1, 2500).reshape((50, 50))
linear1 = np.linspace(0, 255, 2500).reshape((50, 50)).astype(np.uint8)

print("Linear0:", linear0.dtype, linear0.min(), linear0.max())
print("Linear1:", linear1.dtype, linear1.min(), linear1.max())

fig, (ax0, ax1) = plt.subplots(1, 2, figsize=(15, 15))
ax0.imshow(linear0, cmap='gray')
ax1.imshow(linear1, cmap='gray');
```

Linear0: float64 0.0 1.0

Linear1: uint8 0 255



The library is designed in such a way that any data-type is allowed as input, as long as the range is correct (0-1 for floating point images, 0-255 for unsigned bytes, 0-65535 for unsigned 16-bit integers).

You can convert images between different representations by using `img_as_float`, `img_as_ubyte`, etc.:

```
[10]: from skimage import img_as_float, img_as_ubyte

image = data.chelsea()

image_ubyte = img_as_ubyte(image)
image_float = img_as_float(image)

print("type, min, max:", image_ubyte.dtype, image_ubyte.min(), image_ubyte.
      ↪max())
print("type, min, max:", image_float.dtype, image_float.min(), image_float.
      ↪max())
print()
print("231/255 =", 231/255.)
```

```
type, min, max: uint8 0 231
```

```
type, min, max: float64 0.0 0.9058823529411765
```

```
231/255 = 0.9058823529411765
```

Your code would then typically look like this:

```
def my_function(any_image):
    float_image = img_as_float(any_image)
    # Proceed, knowing image is in [0, 1]
```



We recommend using the floating point representation, given that `scikit-image` mostly uses that format internally.

### 1.3 Image I/O

Mostly, we won't be using input images from the `scikit-image` example data sets. Those images are typically stored in JPEG or PNG format. Since `scikit-image` operates on NumPy arrays, *any* image reader library that provides arrays will do. Options include `imageio`, `matplotlib`, `pillow`, etc.

`scikit-image` conveniently wraps many of these in the `io` submodule, and will use whichever of the libraries mentioned above are installed:


```
[11]: from skimage import io

image = io.imread('data/balloon.jpg')

print(type(image))
print(image.dtype)
print(image.shape)
print(image.min(), image.max())

plt.imshow(image);
```

```
<class 'numpy.ndarray'>
uint8
(300, 225, 3)
0 255
```





We also have the ability to load multiple images, or multi-layer TIFF images:

```
[12]: from skimage import data_dir
```

```
!ls $data_dir
```

```
__init__.py          grass.png
__pycache__          gray_morph_output.npz
_binary_blobs.py     green_palette.png
_blobs_3d_fiji_skeleton.tif horse.png
_detect.py           hubble_deep_field.jpg
astronaut.png        ihc.png
astronaut_GRAY_hog_L1.npy lbpcascade_frontalface_opencv.xml
astronaut_GRAY_hog_L2-Hys.npy lfw_subset.npy
block.png            logo.png
brick.png            microaneurysms.png
bw_text.png          moon.png
bw_text_skeleton.npy motorcycle_disp.npz
camera.png           motorcycle_left.png
cells_qpi.npz         motorcycle_right.png
cells_qpi_zipped.zip mssim_matlab_output.npz
checker_bilevel.png  multi.fits
chelsea.png          multipage.tif
chessboard_GRAY.png  multipage_rgb.tif
chessboard_GRAY_U16.tif no_time_for_that_tiny.gif
chessboard_GRAY_U16B.tif orb_descriptor_positions.txt
chessboard_GRAY_U8.npy page.png
chessboard_GRAY_U8.npz palette_color.png
chessboard_RGB.png   palette_gray.png
chessboard_RGB_U8.npy phantom.png
chessboard_RGB_U8.npz rank_filter_tests.npz
clock_motion.png     retina.jpg
coffee.png           rocket.jpg
coins.png             rough-wall.png
color.png             simple.fits
diamond-matlab-output.npz tests
disk-matlab-output.npz text.png
foo3x5x4indexed.png  truncated.jpg
```

```
[13]: ic = io.ImageCollection(data_dir + '/*.png')
```

```
print('Type:', type(ic))
```

```
ic.files
```

```
Type: <class 'skimage.io.collection.ImageCollection'>
```

```
[13]: ['/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/astronaut.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/block.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/brick.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/bw_text.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/camera.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/checker_bilevel.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/chelsea.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/chessboard_GRAY.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/chessboard_RGB.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/clock_motion.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/coffee.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/coins.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/color.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/foo3x5x4indexed.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/grass.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/green_palette.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/horse.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/ihc.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/logo.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/microaneurysms.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/moon.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/motorcycle_left.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/motorcycle_right.png',
'/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
```

```

packages/skimage/data/page.png',
    '/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/palette_color.png',
    '/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/palette_gray.png',
    '/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/phantom.png',
    '/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/rough-wall.png',
    '/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/skimage/data/text.png']

```

```

[14]: import os

f, axes = plt.subplots(nrows=5, ncols=len(ic) // 5 + 1, figsize=(15, 10))

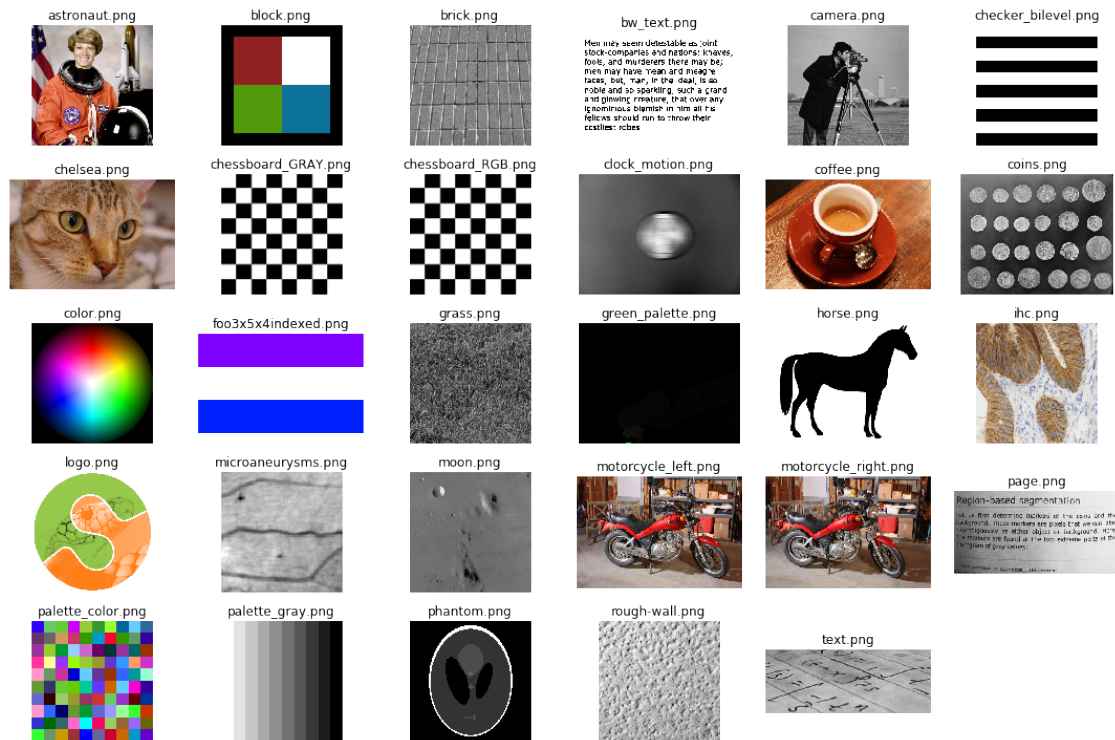
# subplots returns the figure and an array of axes
# we use `axes.ravel()` to turn these into a list
axes = axes.ravel()

for ax in axes:
    ax.axis('off')

for i, image in enumerate(ic):
    axes[i].imshow(image, cmap='gray')
    axes[i].set_title(os.path.basename(ic.files[i]))

plt.tight_layout()

```



## 1.4 Exercise: draw the letter H

Define a function that takes as input an RGB image and a pair of coordinates (row, column), and returns a copy with a green letter H overlaid at those coordinates. The coordinates point to the top-left corner of the H.

The arms and strut of the H should have a width of 3 pixels, and the H itself should have a height of 24 pixels and width of 20 pixels.

Start with the following template:


```
[15]: def draw_H(image, coords, color=(0, 255, 0)):
    out = image.copy()

    ...

    return out
```

Test your function like so:

```
[16]: cat = data.chelsea()
cat_H = draw_H(cat, (50, -50))
plt.imshow(cat_H);
```



## 1.5 Exercise: visualizing RGB channels

Display the different color channels of the image along (each as a gray-scale image). Start with the following template:

```
[17]: # --- read in the image ---  
  
image = plt.imread('data/Bells-Beach.jpg')  
  
# --- assign each color channel to a different variable ---  
  
r = ...  
g = ...  
b = ...  
  
# --- display the image and r, g, b channels ---  
  
f, axes = plt.subplots(1, 4, figsize=(16, 5))  
  
for ax in axes:  
    ax.axis('off')  
  
(ax_r, ax_g, ax_b, ax_color) = axes  
  
ax_r.imshow(r, cmap='gray')
```

```

ax_r.set_title('red channel')

ax_g.imshow(g, cmap='gray')
ax_g.set_title('green channel')

ax_b.imshow(b, cmap='gray')
ax_b.set_title('blue channel')

# --- Here, we stack the R, G, and B layers again
#     to form a color image ---
ax_color.imshow(np.stack([r, g, b], axis=2))
ax_color.set_title('all channels');

```

```

↳ -----

TypeError                                Traceback (most recent call last)

<ipython-input-17-e057ff15c311> in <module>
    18 (ax_r, ax_g, ax_b, ax_color) = axes
    19
--> 20 ax_r.imshow(r, cmap='gray')
    21 ax_r.set_title('red channel')
    22

~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/
↳ __init__.py in inner(ax, data, *args, **kwargs)
    1599     def inner(ax, *args, data=None, **kwargs):
    1600         if data is None:
-> 1601             return func(ax, *map(sanitize_sequence, args), **kwargs)
    1602
    1603         bound = new_sig.bind(ax, *args, **kwargs)

~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/cbook/
↳ deprecation.py in wrapper(*args, **kwargs)
    367         f"%(removal)s. If any parameter follows {name!r}, they "
    368         f"should be pass as keyword, not positionally.")
--> 369     return func(*args, **kwargs)
    370
    371     return wrapper

~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/cbook/
↳ deprecation.py in wrapper(*args, **kwargs)

```

```

367             f"%(removal)s. If any parameter follows {name!r}, they "
368             f"should be pass as keyword, not positionally.")
--> 369         return func(*args, **kwargs)
370
371     return wrapper

~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/axes/
-> _axes.py in imshow(self, X, cmap, norm, aspect, interpolation, alpha, vmin,
-> vmax, origin, extent, shape, filternorm, filterrad, imlim, resample, url,
-> **kwargs)
5669                 resample=resample, **kwargs)
5670
-> 5671         im.set_data(X)
5672         im.set_alpha(alpha)
5673         if im.get_clip_path() is None:

~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/image.
-> py in set_data(self, A)
683             not np.can_cast(self._A.dtype, float, "same_kind")):
684                 raise TypeError("Image data of dtype {} cannot be
-> converted to "
--> 685                             "float".format(self._A.dtype))
686
687         if not (self._A.ndim == 2

```

TypeError: Image data of dtype object cannot be converted to float

Now, take a look at the following R, G, and B channels. How would their combination look? (Write some code to confirm your intuition.)



```
[18]: from skimage import draw

red = np.zeros((300, 300))
green = np.zeros((300, 300))
blue = np.zeros((300, 300))

r, c = draw.circle(100, 100, 100)
red[r, c] = 1

r, c = draw.circle(100, 200, 100)
green[r, c] = 1

r, c = draw.circle(200, 150, 100)
blue[r, c] = 1

f, axes = plt.subplots(1, 3)
for (ax, channel) in zip(axes, [red, green, blue]):
    ax.imshow(channel, cmap='gray')
    ax.axis('off')
```



```
[19]: # Hint: np.stack(..., axis=2)
```

## 1.6 Exercise: Convert to grayscale (“black and white”)

The *relative luminance* of an image is the intensity of light coming from each point. Different colors contribute differently to the luminance: it’s very hard to have a bright, pure blue, for example. So, starting from an RGB image, the luminance is given by:

$$Y = 0.2126R + 0.7152G + 0.0722B$$

Use Python 3.5’s matrix multiplication, `@`, to convert an RGB image to a grayscale luminance image according to the formula above.

Compare your results to that obtained with `skimage.color.rgb2gray`.

Change the coefficients to 1/3 (i.e., take the mean of the red, green, and blue channels, to see how that approach compares with `rgb2gray`).

```
[20]: from skimage import color, img_as_float

image = img_as_float(io.imread('data/balloon.jpg'))

gray = color.rgb2gray(image)
my_gray = ...

# --- display the results ---

f, (ax0, ax1) = plt.subplots(1, 2, figsize=(10, 6))

ax0.imshow(gray, cmap='gray')
ax0.set_title('skimage.color.rgb2gray')

ax1.imshow(my_gray, cmap='gray')
ax1.set_title('my rgb2gray')
```

```
↳ -----

TypeError                                Traceback (most recent call last)

<ipython-input-20-a5fdfebe8302> in <module>
    13 ax0.set_title('skimage.color.rgb2gray')
    14
--> 15 ax1.imshow(my_gray, cmap='gray')
    16 ax1.set_title('my rgb2gray')

~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/
↳ __init__.py in inner(ax, data, *args, **kwargs)
    1599     def inner(ax, *args, data=None, **kwargs):
    1600         if data is None:
-> 1601             return func(ax, *map(sanitize_sequence, args), **kwargs)
    1602
    1603         bound = new_sig.bind(ax, *args, **kwargs)

~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/cbook/
↳ deprecation.py in wrapper(*args, **kwargs)
    367         f"%(removal)s. If any parameter follows {name!r}, they "
    368         f"should be pass as keyword, not positionally.")
--> 369     return func(*args, **kwargs)
    370
```

```

371         return wrapper

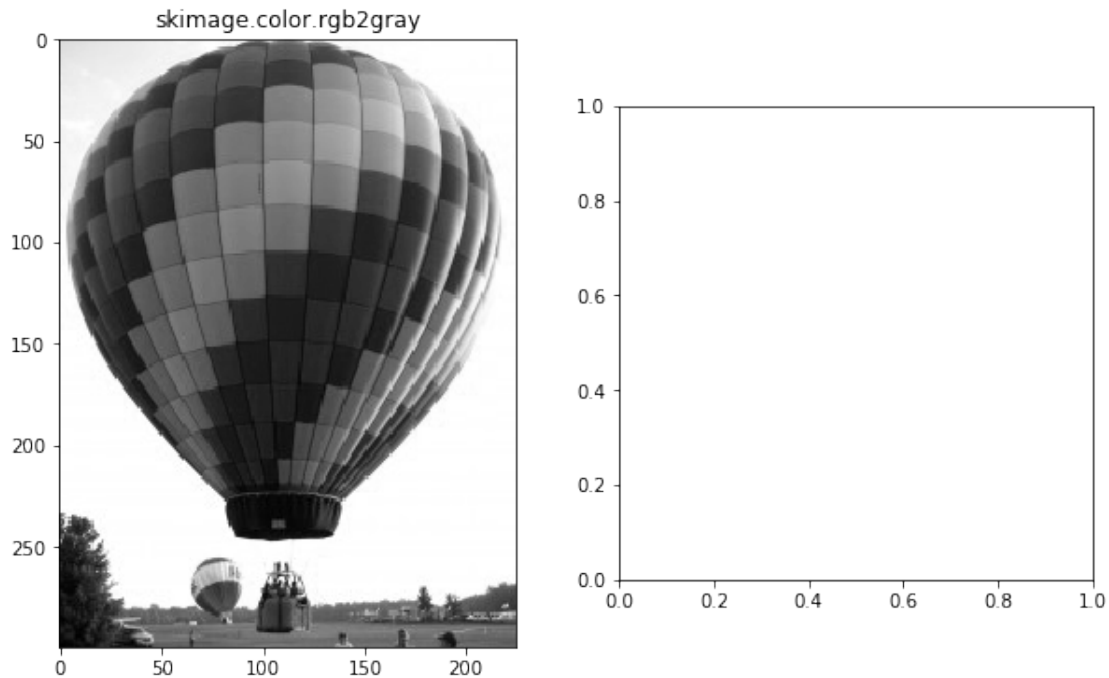
~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/cbook/
↳deprecation.py in wrapper(*args, **kwargs)
    367             f"%(removal)s. If any parameter follows {name!r}, they "
    368             f"should be pass as keyword, not positionally.")
--> 369         return func(*args, **kwargs)
    370
    371         return wrapper

~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/axes/
↳_axes.py in imshow(self, X, cmap, norm, aspect, interpolation, alpha, vmin,
↳vmax, origin, extent, shape, filternorm, filterrad, imlim, resample, url,
↳**kwargs)
    5669                 resample=resample, **kwargs)
    5670
-> 5671         im.set_data(X)
    5672         im.set_alpha(alpha)
    5673         if im.get_clip_path() is None:

~/anaconda3/envs/imagexd19/lib/python3.7/site-packages/matplotlib/image.
↳py in set_data(self, A)
    683             not np.can_cast(self._A.dtype, float, "same_kind")):
    684             raise TypeError("Image data of dtype {} cannot be
↳converted to "
--> 685                             "float".format(self._A.dtype))
    686
    687         if not (self._A.ndim == 2

```

TypeError: Image data of dtype object cannot be converted to float



## 1.7 Just for fun

### 1.7.1 Demo: skimage + keras

```
[25]: from tensorflow.keras.applications.inception_v3 import InceptionV3, \
      ↪ preprocess_input, decode_predictions
      net = InceptionV3()
```

```
/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/tensorflow/python/framework/dtypes.py:526: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
```

```
/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/tensorflow/python/framework/dtypes.py:527: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)]
```

```
/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/tensorflow/python/framework/dtypes.py:528: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)]
```

```
/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
```

```
packages/tensorflow/python/framework/dtypes.py:529: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint16 = np.dtype(["quint16", np.uint16, 1])
/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/tensorflow/python/framework/dtypes.py:530: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint32 = np.dtype(["qint32", np.int32, 1])
/Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/tensorflow/python/framework/dtypes.py:535: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
np_resource = np.dtype(["resource", np.ubyte, 1])
```

```
WARNING:tensorflow:From /Users/dani/anaconda3/envs/imagexd19/lib/python3.7/site-
packages/tensorflow/python/ops/resource_variable_ops.py:435: colocate_with (from
tensorflow.python.framework.ops) is deprecated and will be removed in a future
version.
```

Instructions for updating:

Colocations handled automatically by placer.

Downloading data from [https://github.com/fchollet/deep-learning-models/releases/download/v0.5/inception\\_v3\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels.h5](https://github.com/fchollet/deep-learning-models/releases/download/v0.5/inception_v3_weights_tf_dim_ordering_tf_kernels.h5)  
96116736/96112376 [=====] - 12s 0us/step

```
[26]: from skimage import transform

def inception_predict(image):
    # Rescale image to 299x299, as required by InceptionV3
    image_prep = transform.resize(image, (299, 299, 3), mode='reflect')

    # Scale image values to [-1, 1], as required by InceptionV3
    image_prep = (img_as_float(image_prep) - 0.5) * 2

    predictions = decode_predictions(
        net.predict(image_prep[None, ...])
    )

    plt.imshow(image, cmap='gray')

    for pred in predictions[0]:
        (n, klass, prob) = pred
        print(f'{klass:>15} ({prob:.3f})')
```


```
[27]: from skimage import data, img_as_float
inception_predict(data.chelsea())
```

Downloading data from <https://storage.googleapis.com/download.tensorflow.org/dat>

```
a/imagenet_class_index.json
```


```
40960/35363 [=====] - 0s 0us/step
```

```
    Egyptian_cat (0.904)  
      tabby (0.054)  
    tiger_cat (0.035)  
      lynx (0.000)  
    plastic_bag (0.000)
```



```
[28]: inception_predict(data.camera())
```


```
    tripod (0.937)  
    crutch (0.002)  
  binoculars (0.002)  
reflex_camera (0.001)  
    backpack (0.000)
```



```
[29]: inception_predict(data.coffee())
```

```
    espresso (0.982)  
      cup (0.002)  
  coffee_mug (0.001)  
    eggnog (0.001)  
 espresso_maker (0.001)
```





[ ]: